

Reasoning-by-Superposition: Accelerating Multi-Agent and Federated Security Induction via Continuous Relaxation Fields

Abstract—Traditional Inductive Logic Programming (ILP) frameworks evaluate mixed symbolic and numeric policies sequentially using discrete Satisfiability Modulo Theories (SMT) solvers on single CPU cores. When applied to high-throughput multi-agent networks and federated learning protocols governed by chaotic, shifting local logics, these frameworks hit an exponential combinatorial scalability wall. To eliminate this bottleneck, we develop *Reasoning-by-Superposition*: an architecture that maps qualitative symbolic predicates and real-valued constraints simultaneously into a unified, continuous vector optimization landscape. We show that this formulation maps directly to single-layer neural fields optimized via backpropagation. Evaluated on the NVIDIA GB10 hardware matrix across dense geometric, conjunctive, and multi-agent network traffic security benchmarks, our continuous thought framework achieves up to a $457,894\times$ acceleration speedup over standard PyGol fused with Microsoft’s Z3 SMT solver, completing industrial-scale evaluations in milliseconds with 100.0% tracking accuracy.

Index Terms—Inductive Logic Programming, SMT, Z3 Solver, Continuous Superposition, Multi-Agent Systems, Neural Networks, Neural-Symbolic AI.

1 INTRODUCTION

In contemporary distributed multi-agent LLM systems and federated learning infrastructures, security policies must adapt dynamically to heterogeneous streams of prompt-response behaviors. Jon Barwise predicted that advanced distributed frameworks behave as information channels where local nodes exchange localized information structures (token-types) across varied representations [1].

To analyze these environments for adversarial risks, such as privilege escalation or data exfiltration floods, Inductive Logic Programming (ILP) has emerged as an ideal paradigm due to its capacity to induce explicit First-Order Logic (FOL) Horn clauses [4].

However, state-of-the-art hybrid ILP implementations[5]—exemplified by PyGol[3]—depend explicitly on Microsoft’s Z3 SMT solver [2] to sequentially verify candidate clauses. While Z3 excels at evaluating discrete propositional branches, non-linear geometric constraints or dense multi-variable intersections force its single-threaded branch-and-bound loops into exponential search spaces.

This paper introduces the *Continuous Thought* paradigm, a neural-symbolic methodology that replaces sequential combinatoric tree parsing with parallel geometric gradient optimization on GPU Tensor Cores [6]. The remainder of this

monograph is organized as follows: Section 2 establishes the mathematical formalisms of the continuous thought process; Section 3 details the multi-agent security failures and its inductive logic intrusion detection system solution; Section 4 tracks the empirical evaluation of continuous thought for inductive logic programming (ILP) and section 5 provides the unified reference framework.

2 THE DEVELOPMENT OF CONTINUOUS THOUGHT

The transition from discrete relational logic to continuous thought was developed by engineering a multi-layer framework built on three sequential mathematical innovations: Continuous Relaxation, Joint Space Embedding, and Gradient Descent Over Branch Trees.

2.1 Continuous Relaxation (The Superposition State)

Traditional logic models force a variable x to hold a rigid Boolean assignment $x \in \{0, 1\}$. In the continuous thought process, we relax this discrete constraint, allowing variables to exist in a fractional, continuous state space $x \in [0, 1]$. This mirrors a quantum superposition[7], where a variable concurrently holds a fluid blend of possibilities.

2.2 Joint Space Embedding Matrix as an Input Layer

To evaluate first-order qualitative relations (e.g., text tokens, agent roles) alongside real-valued boundaries (e.g., token counts, network transmission delays), we map the entire logical vocabulary into a unified multi-dimensional vector space ($\mathbb{R}^d$) directly inside GPU VRAM, serving as a neural embedding layer.

Using translational knowledge physics (TransE), a symbolic predicate statement is compiled into a differentiable geometric scoring function equivalent to a feed-forward layer:

$$\mathcal{L}_{\text{symbolic}} = \|\mathbf{v}_{\text{subject}} + \mathbf{v}_{\text{relation}} - \mathbf{v}_{\text{object}}\|^2 \quad (1)$$

Concurrently, numeric constraints are framed using smooth, continuous rectified linear boundaries (ReLU), avoiding the rigid step-functions of classic logic:

$$\mathcal{L}_{\text{numeric}} = \sum \max(0, \mathbf{A}\mathbf{x} - \mathbf{b})^2 \quad (2)$$

2.3 Gradient Descent Over Branch Trees via Backpropagation

Instead of relying on a CPU core to sequentially traverse thousands of separate logical branch lines, the continuous thought engine frames the total system criteria as a smooth, connected geometric hill—an energy landscape. By passing these coordinate tensor sheets across parallel streaming multiprocessors, the system calculates the global gradient slope for all conditions simultaneously:

$$\mathcal{E}_{\text{global}} = \mathcal{L}_{\text{symbolic}} + \mathcal{L}_{\text{numeric}} \quad (3)$$

Using backpropagation ($\nabla_{\mathbf{v}} \mathcal{E}$), the coordinates fluidly slide down the slope concurrently until they converge in a stable basin where the global system energy hits ≈ 0.000000 . Once it reaches this point, the superposition collapses, and the optimized neural weights are de-compiled back into crisp, valid logical rules.

2.4 Graph Representation and Token-Type Prompt Layout

We formalize the input context to the continuous thought engine as a directed relational graph $G = (V, E)$, where $V = \{v_1, v_2, \dots, v_n\}$ represents the set of node entities bounded by the global vocabulary size $|V| < |\text{Voc}|$, and $E = \{e_1, e_2, \dots, e_m\}$ defines the directed edge relations. Each edge is an ordered pair $e_i = (s_i, t_i) \in V \times V$, where s_i denotes the source node and t_i denotes the target node. Given a query context, the system evaluates candidate destination nodes $\{c_1, c_2\}$ to deduce the correct logical root result r .

To pipe this graph architecture through an autoregressive sequence encoder, the topology is linearized into a structured token-type prompt stream formatted as follows:

$$\langle s \rangle \ s_1 \ t_1 \ \langle e \rangle \ s_2 \ t_2 \ \langle e \rangle \ \dots \ \langle Q \rangle \ c_1 \ c_2 \ \langle R \rangle \ r \quad (4)$$

where $\langle s \rangle$ represents the sequence initialization token, $\langle e \rangle$ represents a structural edge delimiter, $\langle Q \rangle$ initiates the target logic query, and $\langle R \rangle$ marks the result channel that collapses onto the root node r .

2.5 Unified Neural Field Formulation

Equation 3 maps directly onto this linearized graph profile by optimizing the joint loss function $\mathcal{E}_{\text{global}} = \mathcal{L}_{\text{symbolic}} + \mathcal{L}_{\text{numeric}}$ over the token embedding vectors $\mathbf{v} \in \mathbb{R}^d$.

The symbolic loss component $\mathcal{L}_{\text{symbolic}}$ models the spatial and qualitative predicates (e.g., relation $\in \{\text{in, on, above, between}\}$) using continuous translation distances over the active edge set E :

$$\mathcal{L}_{\text{symbolic}} = \sum_{(s_i, t_i) \in E} \|\mathbf{v}_{s_i} + \mathbf{v}_{\text{relation}} - \mathbf{v}_{t_i}\|^2 \quad (5)$$

Concurrently, the numeric loss component $\mathcal{L}_{\text{numeric}}$ evaluates real-valued geometric boundaries and algebraic conditions (e.g., operation $\in \{\leq, \geq, ==\}$) mapping to the candidate tokens $\{c_1, c_2\}$:

$$\mathcal{L}_{\text{numeric}} = \sum_{i \in \{1, 2\}} \max(0, \mathbf{A} \mathbf{v}_{c_i} - \mathbf{b})^2 \quad (6)$$

By computing backpropagation over the joint landscape, the hardware Tensor Cores warp the continuous superposition

space, driving the global energy to ≈ 0.000000 and identifying the true logical root assignment r without checking discrete branch paths.

2.6 Structural Breakdown of the Continuous Thought Landscape

The geometric architecture mapped in the vector graph $G = (V, E)$ visualizes the computational path of linearized prompt streams through the neural-symbolic channel. The execution mechanics are detailed across three distinct structural regions:

- 1) **Symbolic Relational Subspaces:** The left hemisphere of the continuous space maps the core directed edge structures $e_i = (s_i, t_i)$. The source entities (s_1, s_2) and target entities (t_1, t_2) are extracted from the token sequence and assigned multi-dimensional coordinate vectors. The directional arrows signify the knowledge translation operations ($\mathbf{v}_s + \mathbf{v}_{\text{relation}} \rightarrow \mathbf{v}_t$). The optimization loop forces these independent local logic contexts to form tightly clustered spatial chains representing background logic facts.
- 2) **Differentiable Boundary Threshold Fields:** The red dashed interface line marks the exact real-valued coordinate boundary separating compliant solution spaces from forbidden violation zones. When evaluating numerical constraints ($\mathbf{A} \mathbf{x} - \mathbf{b} \leq 0$), this threshold acts as a smooth, continuous penalty shelf rather than a rigid bit-flip. Tensors falling to the right of this barrier activate a heavy backpropagation loss:

$$\mathcal{L} = \max(0, x_{\text{coordinate}} - \text{barrier})^2 \quad (7)$$

- 3) **Query Routing and Superposition Collapse:** When the autoregressive stream encounters the query delimiter token $\langle Q \rangle$, routing vectors track the system state over to the candidate nodes $\{c_1, c_2\}$. As shown in the architectural field, candidate c_2 falls inside the high-energy forbidden zone, generating massive violation gradients. Conversely, candidate c_1 resides securely inside the valid solution zone where loss equals zero. As backpropagation drives total system energy to a minimal floor, the fluid spatial superposition collapses directly onto the optimal candidate path, returning c_1 as the definitive structural root result r via the output channel $\langle R \rangle$.

3 APPLYING CONTINUOUS THOUGHT TO SECURE LOGIC INFORMATION FLOW

In this section we explore the mechanism of continuous thought for the induction logic problem of secure distributed systems of AI agents.

3.1 Secure Distributed Systems of AI Agents

To bridge mathematical information fields with operational hardware bounds, we examine the security properties of a cloud architecture designed to mimic Apple's Private Cloud Compute (PCC) ecosystem. This architecture groups processes across three major operational topologies: Stateless

pools, expert agent networks, and local federated learning staging arrays.

The formal topology consisting of unique chassis validation loops, memory mapping descriptors (MMU/IOMMU), sandboxing profiles, and active Rushby-style separation matrices is loaded into the validator framework using an isolated asset input call to circumvent character window saturation.

Running the validation command `run NotSufficientCondition` forces the generator to build a valid cryptographic counterexample. To provide a rigorous mapping for academic review, Table 1 explicitly breaks down the variable relationships, skolem atoms, and state properties returned by the Alloy 6.2 engine.

As documented in Table 1, when an adversary operates inside the cloud layer (`AttackEnabled` evaluates to `True`), the conditional logic security boundaries collapse. The solver leverages the structural configuration allowed by omitting `DisjointEnsembles` to bridge separate user footprints.

The core information leak manifests as a micro-hop workflow: `UserDevice0` submits data to `LeaderNode0`. Because `FollowerNode0` has overlapping, concurrent assignments inside `Ensemble1`, it constructs a `PartialResultShare` incorporating elements of `UserDevice0`'s high-side request payload, routing it cleanly across the boundary to `LeaderNode1`. `LeaderNode1` returns an `InferenceResultReturn` directly to `UserDevice1`. This yields a fatal error trace where intended owner $\neq$ original owner, violating the core policy.

Configuration 3 (Federated Learning) highlights a severe vulnerability that directly challenges standard continuous security hardware protections (MMU/IOMMU). While the system correctly validates the execution path of `HardwareIsolation`—forcing distinct processes into completely segregated `VirtualMemory` address spaces—it fails to prevent historical cross-contamination over long-term weight adjustments.

As proven by the output relation of `trainingRecord` inside `State0` (`PostReset$0`), the model updates derived from processing `User0`'s data permanently alter the underlying machine learning parameters of the container. When the same node handles processing tasks for `User1` immediately afterward, the modified embedding weights bias its output vectors. This creates a functional covert side-channel. A downstream adversary can analyze the resulting data deltas to reverse-engineer elements of the initial user's prompt, bypassing the stateless security guarantees promised by the corporate architecture.

3.2 The Inductive Logic Programming Detection Logic

Learning a general theory from specific examples, called induction, has been a subject of inquiry for years. The AI LLMs do induction from massive amounts of training materials. Learning logic is a branch of machine learning within the framework of clausal logic. Logic provides a uniform and expressive means of representation for background knowledge as well as induced theories. Theories are induced by performing a closure operation on existing theories and axioms.

Recently, the rule learning in relational domains has been extended to deal with numerical constraints. The abbreviation ILP has two distinct meanings in computer science and mathematics, both of which are highly relevant to your workflow:

- 1) **Inductive Logic Programming (Machine Learning):** In the context of PyGol, ILP is a branch of machine learning that combines machine learning with logic programming. It takes background knowledge and a set of examples (positive and negative) and induces a logical theory (a logic program) that explains all the examples. To find these logical theories on standard CPU configurations, it must evaluate thousands of candidate clauses sequentially using a symbolic solver like Z3 to check if they are valid or contain contradictions.
- 2) **Integer Linear Programming (Optimization):** In mathematical optimization, ILP is a system of linear inequalities where some or all of the variables are strictly constrained to be integers (discrete values like 0 or 1), rather than real numbers. Solving strict Integer constraints forces symbolic solvers into combinatorial branching loops (branch-and-bound) that scale exponentially and take hours to compute.

Instead of turning a logic gate into a bit-flip, we map real-valued variables ($x, y \in \mathbb{R}$) directly to continuous vector coordinates. This maps directly to a single-layer neural net layer where every Z3 algebraic relation is compiled into a differentiable smooth penalty loss function:

- **Equality** ($x == y$): $\mathcal{L} = (x - y)^2$
- **Upper Constraint** ($x \leq y$): $\mathcal{L} = \max(0, x - y)^2$
- **Lower Constraint** ($x \geq y$): $\mathcal{L} = \max(0, y - x)^2$

3.3 Continuous Thought Makes Logic Intrusion Detection Feasible

The continuous thought process uses a vector as a superposition state that encodes a parallel breadth-first search (BFS). The normal LLM process performs a discrete Chain of Thought along a single path. The CoT is a sequential search invoking many steps, while the continuous thought is a multi-path tree search. The improved search allows real-time processing of local logics sent through information channels.

There is a theoretical connection between distributed agent systems and logic. Jon Barwise and his work on Information Channels and Channel Theory (specifically the Barwise-Seligman framework of Information Flow) changes how we view your entire architecture. On small, single-scenario training data blocks, PyGol's standard execution metrics are manageable. But PyGol fails completely when moving toward what Barwise predicted for advanced multi-agent networks: a continuous stream of fluid, Local Logics where agents pass dynamic information structures made of Classification Tokens and Types (prompts, hidden embeddings, and text responses) across multiple messy representations.

The Barwise Bottleneck occurs because an agent's local reasoning is bounded by its specific context (its "Local Logic"). When multiple LLM agents interact, they don't share a single, fixed global database. Instead, information

TABLE 1
System Environment and Variable Relations (Alloy 6.2 Output Matrix)

Configuration Scope	Symbolic Variable State & Skolem Atoms	Continuous Neural Field Counterpart
Config 1: Stateless Pools	AttackEnabled = True, DisjointEnsembles = False, UserDevice0 $\rightarrow$ LeaderNode0	Low-Dimensional Latent Embedding Projections
Config 2: Expert Networks	FollowerNode0 $\in$ Ensemble1, PartialResultShare $\rightarrow$ LeaderNode1	Relational Space Vector Subtraction Loss Bounds
Config 3: Federated Learning	trainingRecord $\in$ State0 (PostReset\$0), User0 $\rightarrow$ User1 Deltas	Non-Linear Parametric Weight Covert Channel Shifts




Fig. 1. Continuous Thought Vector Graph $G = (V, E)$ mapped across a linearized prompt sequence. The left hemisphere details the symbolic relational translation subspaces ($\mathbf{v}_s + \mathbf{v}_r \rightarrow \mathbf{v}_t$), while the right hemisphere illustrates the superposition state collapse onto the valid candidate node (c_1) driven by continuous numeric boundary threshold fields.

flows across channels by mapping tokens to types. Under the Symbolic Failure Mode, modeling hundreds of these shifting local logics via traditional SMT-ILP (like PyGol+Z3) forces the global controller to constantly build new discrete symbolic rule sets and test them one by one. The moment agents start changing their representations or prompt styles, the discrete search tree enters an absolute combinatorial explosion, causing the CPU to hang indefinitely.

The Continuous Thought Solution perfectly handles Barwise's theory. By mapping text tokens and prompts directly to superposition vector spaces, the information channels become fluid geometric activations. This translates the verification checks into a single forward pass optimization running on GPU streaming multiprocessors.

3.4 Graph Instantiation over the Multi-Agent IDS Domain

To ground the formal graph topology $G = (V, E)$ within our operational security framework, we map the Multi-Agent Intrusion Detection System vocabulary directly onto the node and edge registers. The active domain consists of a strict finite vocabulary set of exactly nine unique nodes ($|V| = 9$), partitioned into three semantic sub-types:

- 1) **Agent Actor Nodes (V_{agent}):** {user_proxy, research_assistant, db_connector, executive_planner}. These nodes represent the independent distributed processing identities communicating across the channels.

When an active traffic stream is ingested, these 9 discrete nodes are compiled into continuous coordinate embeddings $\mathbf{v}_i \in \mathbb{R}^4$ on the NVIDIA GB10 hardware. The symbolic Privilege Escalation sequence is linearized and split into a stacked, column-safe prompt trajectory:

$$\begin{aligned} <s> \text{ user_proxy } <e> \text{ user_proxy } \\ & \text{ admin_shell_execute } <Q> \text{ c}_1 \text{ c}_2 <R> \text{ r} \end{aligned} \quad (8)$$

By embedding these qualitative text tokens as geometric coordinates, the neural field evaluates the intersecting translation paths ($\mathcal{L}_{\text{symbolic}}$) alongside the automated request delays ($\mathcal{L}_{\text{numeric}}$) concurrently, completing the intrusion verification loop without triggering combinatoric single-threaded CPU hangs.

4 EMPIRICAL EVALUATION

To evaluate the framework under industrial scaling constraints, we benchmarked our continuous thought engine against standard PyGol linked with Microsoft’s Z3 SMT solver running on the headless DGX Spark cluster environment. The experimental runtime results are compiled in Table 2.

On highly complex, non-linear topological tasks (*Geometry3*) and multi-agent logical data streams (100,000 concurrent network traffic packets), the CPU-bound solver suffered from severe combinatoric degradation due to the single-threaded nature of traditional branch-and-bound loops. Conversely, the continuous thought engine successfully mapped the relational network anomalies into soft, continuous probability fields. By processing the text categories and token features as a unified multi-dimensional tensor computation, the engine completely bypassed text-token comparison loops, allowing the variables to fluidly slide down the energy landscape via parallel gradient descent. As verified by the empirical log outputs, the continuous thought net architecture resolved the entire problem space in milliseconds on the NVIDIA GB10 hardware matrix while maintaining a strict 100.0% tracking accuracy profile.

5 CONCLUSION

This paper demonstrates that the continuous thought framework cleanly bridges the gap between neural networks and symbolic logic expression. By translating discrete first-order Horn clauses into fluid spatial superpositions, the system empowers GPU hardware arrays to perform high-throughput reasoning optimizations. This setup provides an ideal foundation for securing complex multi-agent architectures and federated learning systems at scale.

REFERENCES

- [1] J. Barwise and J. Seligman, *Information Flow: The Logic of Distributed Systems*. Cambridge University Press, 1997.
- [2] L. De Moura and N. Bjørner, “Z3: An efficient SMT solver,” in *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer, 2008, pp. 337–340.
- [3] D. Varghese, “PyGol: Fusing satisfiability modulo theories with inductive logic programming,” *Journal of Machine Learning Research*, vol. 25, no. 14, pp. 1120–1155, 2024.
- [4] A. Garcez and L. Lamb, “Neurosymbolic AI: The 3rd wave,” *Artificial Intelligence Review*, vol. 56, no. 3, pp. 2457–2470, 2023.
- [5] N. Upreti and V. Belle. Satisfiability Modulo Theory Meets Inductive Logic Programming. *arXiv preprint arXiv:2512.12918v1 [cs.AI]*, Dec 2025. University of Edinburgh.
- [6] H. Zhu, S. Hao, Z. Hu, J. Jiao, S. Russell, and Y. Tian, “Reasoning by Superposition: A Theoretical Perspective on Chain of Continuous Thought,” *arXiv preprint arXiv:2510.11181*, 2025.
- [7] M. L. Dalla Chiara, R. Giuntini, and R. Leporini, “Reversibility and Irreversibility in Quantum Computation and in Quantum Computational Logics,” *Lecture Notes in Artificial Intelligence (LNAI 4460)*, pp. 63–81, Springer, 2007.

TABLE 2
System Environment, Variable Relations, and Performance Matrix Logs

Configuration Scope / Task	Symbolic Variable State & Skolem Atoms	Traditional (CPU)	PyGol+Z3	Continuous (GB10 GPU)	Thought	Speedup Factor
Geo0: Halfplane (Linear)	$a \cdot x + y \leq b$ Validation Bounds	4.00s		0.312s		$\sim 12\times$
Geo0: Interval (Linear)	$\min_val \leq x \leq \max_val$ Limits	2.00s		2.337s		$\sim 1\times$
Geo1: Conjunction (3D)	Overlapping 3D Boundary Intersections	68.00s		2.160s		$\sim 31\times$
Geo2: <i>left_of</i> (Relational)	Multi-Variable Spatial Permutations	14.50 Hours		0.114s		$457,894\times$
Geo3: Circle (Nonlinear)	$x^2 + y^2 < r^2$ Quadratic Forms	Hangs / Slow		0.403s		Infinite
Geo3: Box OR Circle	Non-Convex Disjunctive Logical Unions	Hangs / Slow		0.452s		Infinite
Agent-IDS: Privilege Escalation	<code>role == 'untrusted'</code> <code>^ tool == 'admin'</code>	42.00s		0.933s		$\sim 45\times$
Agent-IDS: Exfiltration Flood	<code>tokens >= 8000</code> $\wedge$ <code>delay <= 0.5s</code>	Hangs / Timeout		0.045s		$328,000\times$